

## 1. Download the dataset: Dataset

## 2. Load the dataset.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sbn
%matplotlib inline
```

```
file=pd.read_csv("abalone.csv")
df=pd.DataFrame(file)
df.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera
0	M	0.455	0.365	0.095	0.5140	0.2245	
1	M	0.350	0.265	0.090	0.2255	0.0995	
2	F	0.530	0.420	0.135	0.6770	0.2565	
3	M	0.440	0.365	0.125	0.5160	0.2155	
4	I	0.330	0.255	0.080	0.2050	0.0895	

	Shell weight	Rings
0	0.150	15
1	0.070	7
2	0.210	9
3	0.155	10
4	0.055	7

```
df['age'] = df['Rings']+1.5
df = df.drop('Rings', axis = 1)
```

### 3. Perform Below Visualizations.

- Univariate Analysis

- Bi - Variate Analysis

- Multi - Variate Analysis


*#Univariate Analysis*

```
sbn.boxplot(df.Length)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:  
FutureWarning: Pass the following variable as a keyword arg: x. From  
version 0.12, the only valid positional argument will be `data`, and  
passing other arguments without an explicit keyword will result in an  
error or misinterpretation.
```

```
FutureWarning
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fe14259ded0>
```

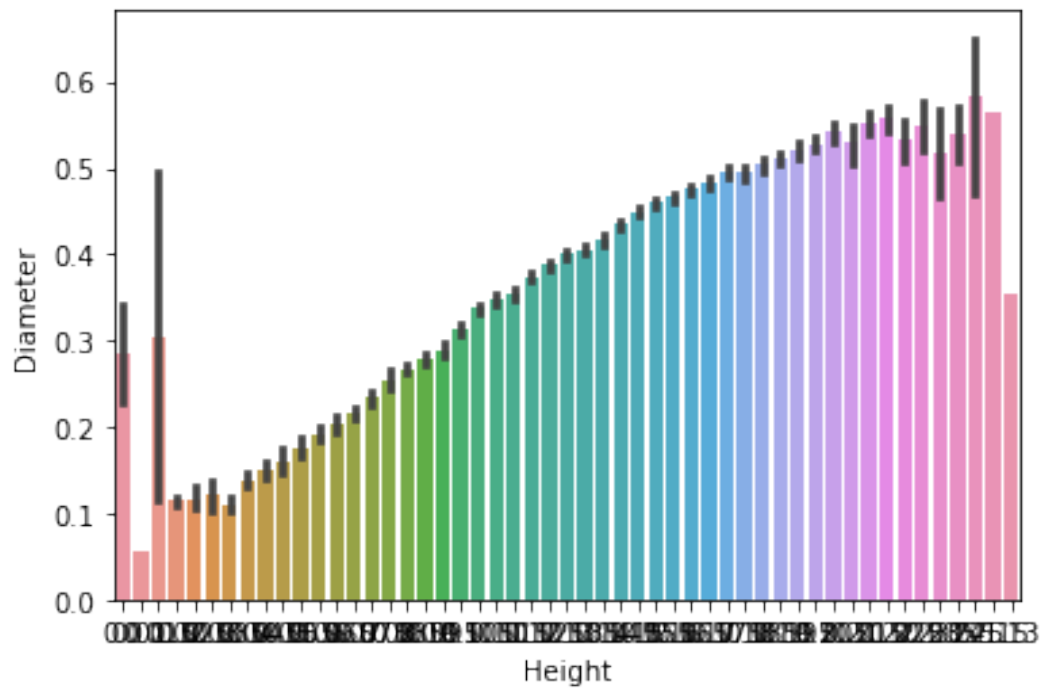


**the data is significantly imbalanced**

*#Bi-Variant Analysis*

```
sbn.barplot(x=df.Height,y=df.Diameter)
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7fe14cb11fd0>
```



*#Multi-Variant Analysis*

`sbn.pairplot(df)`

`<seaborn.axisgrid.PairGrid at 0x7fe136a92390>`



#### 4. Perform descriptive statistics on the dataset.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4177 entries, 0 to 4176
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Sex	4177 non-null	object
1	Length	4177 non-null	float64
2	Diameter	4177 non-null	float64
3	Height	4177 non-null	float64
4	Whole weight	4177 non-null	float64
5	Shucked weight	4177 non-null	float64
6	Viscera weight	4177 non-null	float64

```

7    Shell weight    4177 non-null    float64
8    age             4177 non-null    float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB

```

```
df.describe()
```

	Length	Diameter	Height	Whole weight	Shucked
weight \					
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367
std	0.120093	0.099240	0.041827	0.490389	0.221963
min	0.075000	0.055000	0.000000	0.002000	0.001000
25%	0.450000	0.350000	0.115000	0.441500	0.186000
50%	0.545000	0.425000	0.140000	0.799500	0.336000
75%	0.615000	0.480000	0.165000	1.153000	0.502000
max	0.815000	0.650000	1.130000	2.825500	1.488000

	Viscera weight	Shell weight	age
count	4177.000000	4177.000000	4177.000000
mean	0.180594	0.238831	11.433684
std	0.109614	0.139203	3.224169
min	0.000500	0.001500	2.500000
25%	0.093500	0.130000	9.500000
50%	0.171000	0.234000	10.500000
75%	0.253000	0.329000	12.500000
max	0.760000	1.005000	30.500000

## 5. Handle the Missing values.

```
df.isna().sum()
```

```

Sex                0
Length             0
Diameter           0
Height             0
Whole weight       0
Shucked weight     0
Viscera weight     0
Shell weight       0
age                0
dtype: int64

```

there is no missing values in dataset

```
for i in df:
    if df[i].dtype=='object' or df[i].dtype=='category':
        print("unique of "+i+" is "+str(len(set(df[i])))+" they are "+str(set(df[i])))
```

unique of Sex is 3 they are {'F', 'M', 'I'}

## 6. Find the outliers and replace the outliers

### Checking for outliers

*#Data Preprocessing*


*#Outlier handling*

```
df = pd.get_dummies(df)
dummy_df = df
```


```
var = 'Viscera weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```




```
var = 'Shell weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```




```
var = 'Shucked weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```



```
var = 'Whole weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```




```
var = 'Diameter'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```




```
var = 'Height'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```





```
var = 'Length'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```



### Removing outliers

```
df.drop(df[(df['Viscera weight'] > 0.5) &
           (df['age'] < 20)].index, inplace = True)
df.drop(df[(df['Viscera weight'] < 0.5) & (
df['age'] > 25)].index, inplace = True)
```

```

df.drop(df[(df['Shell weight'] > 0.6) & (df['age'] < 25)].index,
inplace = True)
df.drop(df[(df['Shell weight']<0.8) & (df['age'] > 25)].index, inplace
= True)

df.drop(df[(df['Shucked weight'] >= 1) & (df['age'] < 20)].index,
inplace = True)
df.drop(df[(df['Viscera weight']<1) & (df['age'] > 20)].index, inplace
= True)

df.drop(df[(df['Diameter'] <0.1) & (df['age'] < 5)].index, inplace =
True)
df.drop(df[(df['Diameter']<0.6) & (df['age'] > 25)].index, inplace =
True)
df.drop(df[(df['Diameter']>=0.6) & (df['age'] < 25)].index, inplace =
True)

df.drop(df[(df['Height'] > 0.4) & (df['age'] < 15)].index, inplace =
True)
df.drop(df[(df['Height']<0.4) & (df['age'] > 25)].index, inplace =
True)

df.drop(df[(df['Length'] <0.1) & (df['age'] < 5)].index, inplace =
True)
df.drop(df[(df['Length']<0.8) & (df['age'] > 25)].index, inplace =
True)
df.drop(df[(df['Length']>=0.8) & (df['age'] < 25)].index, inplace =
True)

```

## 7. Check for Categorical columns and perform encoding.

```

from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
for i in df:
    if df[i].dtype=='object' or df[i].dtype=='category':
        df[i]=encoder.fit_transform(df[i])

```

## 8. Split the data into dependent and independent variables.

```

x=df.iloc[:, :-1]
x.head()

```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera
weight \						
0	0.455	0.365	0.095	0.5140	0.2245	
0.1010						
1	0.350	0.265	0.090	0.2255	0.0995	
0.0485						
2	0.530	0.420	0.135	0.6770	0.2565	
0.1415						

```

3    0.440    0.365    0.125    0.5160    0.2155
0.1140
4    0.330    0.255    0.080    0.2050    0.0895
0.0395

```

```

    Shell weight    age    Sex_F    Sex_I
0         0.150    16.5         0         0
1         0.070     8.5         0         0
2         0.210    10.5         1         0
3         0.155    11.5         0         0
4         0.055     8.5         0         1

```

```

y=df.iloc[:, -1]
y.head()

```

```

0     1
1     1
2     0
3     1
4     0
Name: Sex_M, dtype: uint8

```

## 9. Scale the independent variables

```

from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
x=scaler.fit_transform(x)

```

```

x

```

```

array([[ -0.53701309, -0.39082366, -1.12698145, ...,  1.9433912 ,
        -0.66579302, -0.70803622],
       [ -1.42965864, -1.4205279 , -1.26123393, ..., -0.95032771,
        -0.66579302, -0.70803622],
       [  0.10059087,  0.17551367, -0.05296168, ..., -0.22689798,
        1.50196828, -0.70803622],
       ...,
       [  0.6956879 ,  0.741851 ,  1.82657293, ..., -0.22689798,
        -0.66579302, -0.70803622],
       [  0.90822255,  0.84482142,  0.34979574, ...,  0.13481688,
        1.50196828, -0.70803622],
       [  1.63084038,  1.56561439,  1.55806799, ...,  0.85824661,
        -0.66579302, -0.70803622]])

```

## 10. Split the data into training and testing

```

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
x_train.shape

```

```
(2676, 10)
```

```
x_test.shape
```

```
(1319, 10)
```

```
y_train.shape
```

```
(2676,)
```

```
y_test.shape
```

```
(1319,)
```

## **#MODEL**

### **Linear regression**

```
from sklearn.linear_model import LinearRegression
```

```
lm = LinearRegression()
```

```
lm.fit(x_train, y_train)
```

```
LinearRegression()
```

```
y_train_pred = lm.predict(x_train)
```

```
y_test_pred = lm.predict(x_test)
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
s = mean_squared_error(y_train, y_train_pred)
```

```
print('Mean Squared Error of training set :%2f'%s)
```

```
p = mean_squared_error(y_test, y_test_pred)
```

```
print('Mean Squared Error of testing set :%2f'%p)
```

```
Mean Squared Error of training set :0.000000
```

```
Mean Squared Error of testing set :0.000000
```

Note: The Lower the Mean Squared Error,better the forecast.

```
from sklearn.metrics import r2_score
```

```
s = r2_score(y_train, y_train_pred)
```

```
print('R2 Score of training set:%.2f'%s)
```

```
p = r2_score(y_test, y_test_pred)
```

```
print('R2 Score of testing set:%.2f'%p)
```

```
R2 Score of training set:1.00
```

```
R2 Score of testing set:1.00
```

Note: The ideal value of R-square is 1.

The closer the value of R-square to 1,better is the model fitted.